

Operating systems – Assignment 3

I/O Scheduling

Lennart Jern
CS: ens16ljn

Teacher
Ahmed Aley

January 7, 2017

1 Introduction

The Linux kernel provides a number of different scheduling policies that can be used to fine tune the performance of certain applications. In this report, five different schedulers are evaluated using an artificial, CPU intensive, task. Three of the tested schedulers are “normal”, while the last two are “real-time” schedulers, meaning that they provide higher priority for their processes than the normal ones do.

The work load consists of a simple program, called `work`, that sums over a part of Grandi’s series¹ ($1 - 1 + 1 - 1 + \dots$), using a specified number of threads. The source code for `work` can be found in listing ???. Since the task is easy to parallelize, only require minimal memory access and no disk access, it should be comparable to CPU intense tasks like compression and matrix calculations.

2 Method

A Bash script (`timer.sh`, listing 1) was used to time the complete task 10 times for each scheduler, for thread counts ranging from 1 to 10. See code listing 1 for the code. Additionally, each thread keeps track of the time taken from the start of its execution until it is finished, and prints this information to `stdout`, which is forwarded to data files by the bash script.

All this data was then processed by a simple Python program, `stats`, in order to calculate the median, minimum and maximum run time for each scheduler and thread count; for both the total run times and the individual thread times. This program can be found in listing 2. In addition to the statistical calculations, `stats` also produces some figures to describe the data. The figures and calculations are mostly done using the python libraries Pandas² and Matplotlib³.

It should be noted here that the processes running with real-time schedulers were run with maximum priority. Since the other schedulers does not accept a priority setting, other than the nice value, these were left untouched.

All tests were run on my personal computer with the specifications seen in table 1.

¹https://en.wikipedia.org/wiki/Grandi's_series

²<http://pandas.pydata.org/>

³<http://matplotlib.org/>

Component	Specification
OS:	Fedora 25
Kernel:	Linux 4.8.13-300.fc25.x86_64
CPU:	Intel Core i5-2500K CPU @ 3.7GHz
RAM:	7965MiB
GCC:	6.2.1
Bash:	4.3.43
Python:	3.5.2
Pandas:	0.18.1
Matplotlib:	1.5.3

Table 1: Test system specification.

3 Results

	cfq	deadline	noop
count	1.709e+05	1.709e+05	1.709e+05
mean	4.502e-04	4.424e-04	4.408e-04
std	1.485e-03	1.364e-03	1.367e-03
min	2.300e-08	2.300e-08	2.300e-08
25%	8.600e-08	8.300e-08	8.400e-08
50%	2.150e-07	2.160e-07	2.170e-07
75%	1.879e-05	1.866e-05	1.849e-05
max	7.179e-02	2.453e-02	2.202e-02

4 Final thoughts and lessons learned

It is quite clear that it is possible to spend a considerable amount of time just analyzing schedulers. The results are intriguing as they show clear differences in the behavior between the schedulers, and with several different tasks to compare, even more patterns would surely emerge.

A lesson learned is that one should think carefully about when and where to start and stop the timers. If measuring just the whole task, this is quite easy, but for the individual threads it gets more tricky. The thread that starts the other threads must also start the timers since the working thread may not get to start immediately. Similarly, the working threads must stop their own timers, since there might be a pause between the threads finishing and actually getting joined.

A Code listings

Listing 1: timer.sh

```
1  #!/bin/bash
2
3  # timer.sh
4  #
5  # A timer script to measure the differences between i/o schedulers
6  #
7  # Author: Lennart Jern (ens16ljn@cs.umu.se)
8
9  # Clean up old results
10 rm ../data/*.log
11
12 SCHEDULERS="cfq_noop_deadline"
13 DEVICE="/sys/block/sdc/queue/scheduler"
14 # Starting directory and expression to search for
15 START="/run/media/lennart/KINGSTON/test_files_expression"
16 # START="/run/media/lennart/Verbatim/test_files_expression"
17 MNT="/run/media/lennart/KINGSTON"
18
19 for S in $SCHEDULERS
20 do
21     echo $S | sudo tee $DEVICE
22     echo "Scheduler: `cat $DEVICE`"
23     # LINE=""
24     # Time the commands 10 times
25     for i in $(seq 1 10)
26     do
27         # Unmount and mount to clear all cache
28         sudo umount $MNT
29         sudo rm -d /run/media/lennart/KINGSTON
30         sudo mkdir /run/media/lennart/KINGSTON
31         sudo mount /dev/sdc1 $MNT
32
33         # We use 4 parallel commands that store their respective times in
34         # separate log files
35         COMMAND1="./mfind_$START >> ../data/$S-1.log"
36         COMMAND2="./mfind_$START >> ../data/$S-2.log"
37         COMMAND3="./mfind_$START >> ../data/$S-3.log"
38         COMMAND4="./mfind_$START >> ../data/$S-4.log"
39
40         eval $COMMAND1 &
41         eval $COMMAND2 &
42         eval $COMMAND3 &
43         eval $COMMAND4 &
44         # Wait for all commands to finish
45         wait
46
47         # A little progress report
48         echo "Run_$i_done."
49     done
50
51 done
52
53 # Restore cfq scheduler
54 echo "cfq" | sudo tee $DEVICE
```

Listing 2: stats.py

```

1  """
2  stats.py
3
4  Process the data produced by timer.sh by calculating the
5  medians, max values and min values for each scheduler.
6  Also plots the density curves.
7
8  Author: Lennart Jern (ens16ljn@cs.umu.se)
9  """
10
11 import pandas as pd
12 import re
13 import matplotlib.pyplot as plt
14
15 def produce_stats():
16     """
17     Read the data from files, calculate statistical values and make a plot.
18     """
19     base = "../data/"
20     ext = ".csv"
21     header=("cfq", "deadline", "noop")
22
23     # Get individual read times as a DataFrame
24     df = get_read_times()
25
26     stats = df.describe()
27     # Escape per cent chars
28     idx = ['count', 'mean', 'std', 'min', '25%', '50%', '75%', 'max']
29     stats = stats.set_index(idx)
30
31     # Save stats as csv
32     stats.to_csv(base+"stats.csv", header=header, float_format="%.3e")
33
34     # Plot and save the density curves
35     ax = df.plot.kde()
36     ax.set_xlabel("Timer(s)")
37     ax.set_xlim([0, 0.006])
38     fig = ax.get_figure()
39     fig.savefig(base+"density.pdf")
40
41
42 def collect_read_times(file_name):
43     """Read thread times from a file."""
44     f = open(file_name)
45     times = []
46     # Regular expression to find floats
47     time = re.compile("(\\d+\\.\\d+)")
48
49     for line in f:
50         match = time.match(line)
51
52         if (match):
53             t = float(match.group(1))
54             times.append(t)
55
56     return times
57
58 def get_read_times():
59     """Collect timing information about all schedulers in a DataFrame."""
60     schedulers = ["cfq", "deadline", "noop"]
61     base = "../data/"

```

```

62     ext = ".log"
63     header=("cfq", "deadline", "noop")
64
65     # Collect all times in one file
66     times = {key: [] for key in schedulers}
67     for s in schedulers:
68         # We have 4 parallel log files for each scheduler
69         for i in [1,2,3,4]:
70             f = base+s+"-"+str(i)+ext
71             times[s].extend(collect_read_times(f))
72     # Return a DataFrame with all timing data
73     df = pd.DataFrame(times)
74     return df
75
76 # Collect statistical data
77 produce_stats()

```

Listing 3: generate_test_files.sh

```

1  #!/bin/bash
2
3  # File: generate_test_files.sh
4  # Author: Lennart Jern - ens16ljn@cs.umu.se
5  #
6  # Generate a file tree to do tests on.
7
8  # Starting directory
9  # START=/media/removable/KINGSTON
10 START=/run/media/lennart/Verbatim
11
12 # Move to correct directory/device
13 cd $START
14
15 # Remove tree if existent
16 rm -r test_files
17
18 # Create directory to hold all test files
19 mkdir test_files
20 cd test_files
21
22 # Split up the files between a few directories
23 for DIR in a b c d e; do
24     mkdir $DIR
25     cd $DIR
26     # Create empty files
27     for F in $(seq 1 20); do
28         touch $F
29     done
30
31     for D in f g h i j; do
32         mkdir $D
33         cd $D
34     done
35
36     cd "$START/test_files"
37 done
38
39 # Big files
40 mkdir bigs
41 cd bigs
42

```

```

43 # Generate files with lots of zeros...
44 for F in $(seq 1 5); do
45     head -c 50M < /dev/zero > "file$F"
46 done
47
48 cd "$START/test_files"
49 # Generate deep folders
50 for i in $(seq 1 10); do
51     mkdir "deep$i"
52     cd "deep$i"
53     for D in $(seq 1 100); do
54         mkdir "dir$D"
55         cd "dir$D"
56     done
57     cd "$START/test_files"
58 done

```

Listing 4: mfind.c

```

1  /**
2   * File: mfind.c
3   * Author: Lennart Jern - ens16ljn@cs.umu.se
4   *
5   * Usage: ./mfind [-t {d|f|l}] [-p nrthr] start1 [start2 ...] name
6   *
7   * mfind can search after files, links and directories from given start paths.
8   * The search can be done with more than one thread by specifying the flag
9   * '-p#', where # is the number of threads to use.
10  *
11  */
12 #define _GNU_SOURCE
13
14 #include <stdio.h>
15 #include <stdlib.h>
16 #include <errno.h>
17 #include <string.h>
18 #include <dirent.h>
19 #include <sys/stat.h>
20 #include <time.h>           // timing
21 #include "parser.h"         // Includes list.h
22
23 #define ONE_OVER_BILLION 1E-9
24
25 void *find_file(void *s_data);
26 int search_path(SearchData *data, char *path);
27 int search_directory(SearchData *search_data, DIR *dir, char *path);
28 int get_dirent(struct dirent *priv_dirent, DIR *dir);
29 void process_file(char *file_path, char *name,
30                  struct stat f_stat, SearchData *data);
31 int add_dir(LinkedList *list, char *dir_path);
32 void check_starting_dirs(SearchData *search_data);
33 void print_path(void *path);
34 void delete_path(void *path);
35 void SearchData_delete(SearchData *s_data);
36
37 /**
38  * main - parse arguments, do the search and then clean up.
39  * @param argc -- number of arguments
40  * @param argv -- array of arguments
41  * @return      0 if everything went well, a positive int otherwise
42  */

```

```

43 int main(int argc, char *argv[]) {
44     int ret = 0;
45
46     SearchData *search_data = parse_arguments(argc, argv);
47     check_starting_dirs(search_data);
48
49     #ifdef DEBUG
50         printf("#_Search_data\n=====\n");
51         printf("#_Threads:_%d\n", search_data->num_threads);
52         printf("#_Type:_%c\n", search_data->type);
53         printf("#_Needle:_%s\n", search_data->needle);
54         List_print(search_data->directories, print_path);
55         printf("=====\n\n");
56     #endif
57
58     search_data->num_searchers = 0;
59
60     find_file(search_data);
61
62     // Check for errors
63     ret = search_data->error;
64     // Free allocated memory
65     SearchData_delete(search_data);
66     return ret;
67 }
68
69 /**
70  * find_file - search for files and directories
71  * @param s_data -- search data, containing needle to look for and list of
72  *                directories to look in
73  */
74 void *find_file(void *search_data) {
75     unsigned int reads = 0;
76     SearchData *data = search_data;
77     char *path = NULL;
78     int error = 0;
79
80     // Keep searching while there are dirs in the list.
81     while((path = (char *)List_get(data->directories)) != NULL) {
82
83         data->num_searchers++;
84
85         reads++;
86         if (search_path(data, path) != 0) {
87             perror(path);
88             // We don't consider permission denied or missing dir as errors
89             // error = 1;
90         }
91
92         delete_path(path);
93         data->num_searchers--;
94     } // End while. No more dirs to search and all threads done.
95     // Make sure caller knows if there were errors.
96     data->error = error;
97     printf("Reads:_%d\n", reads);
98     return NULL;
99 }
100
101 /**
102  * search_path - open and search the directory given by path
103  * @param data -- SearchData (what to search for)
104  * @param path -- path to directory to search

```



```

105  * @return      0 on successful search, -1 if there were errors
106  */
107  int search_path(SearchData *data, char *path) {
108      // Open the directory. If it fails, clean up and continue with the next one.
109      DIR *dir = opendir(path);
110      int ret = 0;
111
112      if (dir == NULL) {
113          ret = -1;
114          return ret;
115      }
116
117      // Check for matches in the dir
118      if (search_directory(data, dir, path) != 0) {
119          ret = -1;
120      }
121
122      if (closedir(dir) != 0) {
123          perror("closedir");
124      }
125      return ret;
126  }
127
128  /**
129   * search_directory - check all files and folders in dir for matches and
130   * add folders to the list.
131   * @param search_data -- data regarding the search
132   * @param dir          -- dir to look in
133   * @param path         -- path to the dir (used for printing)
134   * @return             0 if everything went well, a poitive int otherwise.
135   */
136  int search_directory(SearchData *search_data, DIR *dir, char *path) {
137      struct dirent *priv_dirent;
138      struct stat f_stat;
139      char *file_path = NULL;
140      int at_end = 0;
141      int ret = 0;
142
143      while (at_end != 1) {
144          priv_dirent = malloc(sizeof(struct dirent));
145          if (priv_dirent == NULL) {
146              perror("malloc");
147              exit(EXIT_FAILURE);
148          }
149
150          at_end = get_dirent(priv_dirent, dir);
151          if (at_end != 0) {
152              // Either error or end of dir
153              if (at_end == -1) {
154                  ret++;
155              }
156              free(priv_dirent);
157              continue;
158          }
159
160          // Build file path string
161          if (asprintf(&file_path, "%s/%s", path, priv_dirent->d_name) == -1) {
162              fprintf(stderr, "Error: asprintf failed. Unable to set file_path.\n");
163              free(priv_dirent);
164              ret++;
165              continue;
166          }

```

```

167     // Get stats (file type)
168     if (lstat(file_path, &f_stat) != 0) {
169         perror(file_path);
170         free(priv_dirent);
171         free(file_path);
172         ret++;
173         continue;
174     }
175
176     // Print matches.
177     process_file(file_path, priv_dirent->d_name, f_stat, search_data);
178
179     // Add directories to the list (not . and ..)
180     if (S_ISDIR(f_stat.st_mode)
181         && strcmp(priv_dirent->d_name, ".") != 0
182         && strcmp(priv_dirent->d_name, "..") != 0) {
183
184         if (add_dir(search_data->directories, file_path) != 1) {
185             fprintf(stderr, "Failed to add directory to list.\n");
186             return -1;
187         }
188     }
189
190     free(file_path);
191     file_path = NULL;
192     free(priv_dirent);
193     priv_dirent = NULL;
194 }
195 return ret;
196 }
197
198 /**
199  * get_dirent - copy the next dirent in dir to priv_dirent in a thread safe way.
200  * This private dirent is safe to use in a multi thread environment.
201  * @param priv_dirent -- pointer to dirent where the dirent will be saved.
202  * @param dir -- dir to read from
203  * @return -1 on error, 1 when the last element was read and
204  *         0 otherwise
205  */
206 int get_dirent(struct dirent *priv_dirent, DIR *dir) {
207     struct dirent *dirent;
208     errno = 0;
209
210     // Starting time
211     struct timespec start;
212     // Time when finished
213     struct timespec end;
214     clock_gettime(CLOCK_REALTIME, &start);
215
216     dirent = readdir(dir);
217
218     // Get the time when finished
219     clock_gettime(CLOCK_REALTIME, &end);
220     // Calculate time it took
221     double time_taken = (end.tv_sec - start.tv_sec)
222         + (end.tv_nsec - start.tv_nsec)
223         * ONE_OVER_BILLION;
224     printf("%.12lf\n", time_taken);
225
226     if (errno != 0) {
227         perror("readdir");
228         return -1;

```

```

229     } else if (dirent == NULL) {
230         // No more files to read
231         return 1;
232     }
233     // Copy dirent to private memory
234     memcpy(priv_dirent, dirent, sizeof(struct dirent));
235     return 0;
236 }
237
238 /**
239  * process_file - print out matching file.
240  * @param file_path -- the path to the file
241  * @param name       -- name of the file
242  * @param f_stat     -- file stats
243  * @param data       -- SearchData (what type and name are we looking for?)
244  */
245 void process_file(char *file_path, char *name,
246                  struct stat f_stat, SearchData *data) {
247     int match = 0;
248     char type = data->type;
249     // Is the name matching?
250     if (strcmp(name, data->needle) == 0) {
251         match = 1;
252     }
253
254     // Check type, print if we have a match
255     if (S_ISDIR(f_stat.st_mode)) {
256         if (match == 1 && (type == 'd' || type == '\0')) {
257             printf("%s\n", file_path);
258         }
259     } else if (S_ISREG(f_stat.st_mode)) {
260         if (match == 1 && (type == 'f' || type == '\0')) {
261             printf("%s\n", file_path);
262         }
263     } else if (S_ISLNK(f_stat.st_mode)) {
264         if (match == 1 && (type == 'l' || type == '\0')) {
265             printf("%s\n", file_path);
266         }
267     }
268 }
269
270 /**
271  * add_dir - add a directory to the list in a thread safe manner
272  * @param list    -- list to add to
273  * @param dir_path -- path to the directory
274  * @return        1 if the dir was added, 0 if addition failed.
275  */
276 int add_dir(LinkedList *list, char *dir_path) {
277     char *new_dir = malloc(strlen(dir_path)+1);
278     if (new_dir == NULL) {
279         perror("malloc");
280         exit(EXIT_FAILURE);
281     }
282     strcpy(new_dir, dir_path);
283
284     if (List_append(list, (void *)new_dir) != 1) {
285         return 0;
286     }
287     return 1;
288 }
289
290 /**

```

```

291  * check_starting_dirs - check if the starting dirs match the search criterias
292  * @param search_data -- data egarding the search
293  */
294  void check_starting_dirs(SearchData *search_data) {
295      char *path;
296      struct stat f_stat;
297      LinkedList *checked_dirs = List_init();
298
299      // Check all starting dirs for matches
300      while((path = (char *)List_get(search_data->directories)) != NULL) {
301          if (lstat(path, &f_stat) != 0) {
302              perror(path);
303              continue;
304          }
305          // Print if there is a match
306          process_file(path, basename(path), f_stat, search_data);
307
308          // Add the checked dir to the new list
309          char *new_dir = malloc(strlen(path)+1);
310          if (new_dir == NULL) {
311              perror("malloc");
312              exit(EXIT_FAILURE);
313          }
314          strcpy(new_dir, path);
315
316          if (List_append(checked_dirs, (void *)new_dir) == 0) {
317              fprintf(stderr, "Could not add path to list.\n");
318              search_data->error++;
319          }
320          free(path);
321      }
322      // Delete the old list
323      List_delete(search_data->directories, delete_path);
324      // Add the checked dirs
325      search_data->directories = checked_dirs;
326  }
327
328  /**
329   * print_path - print out a path
330   * @param path -- a void pointer to a path string
331   */
332  void print_path(void *path) {
333      char *str = (char *)path;
334      printf("%s\n", str);
335  }
336
337  /**
338   * delete_path - delete and free any memory occupied by path
339   * @param path -- path to be freed
340   */
341  void delete_path(void *path) {
342      free(path);
343  }
344
345  /**
346   * searchData_delete - free all memory allocated for s_data
347   * @param s_data -- SearchData to free
348   */
349  void searchData_delete(SearchData *s_data) {
350      free(s_data->needle);
351      List_delete(s_data->directories, delete_path);
352      free(s_data);

```

Listing 5: Makefile

```
1 SOURCE=mfind.c list.c
2 OBJECTS=mfind.o list.o parser.o
3 FLAGS=-std=c11 -Wall -pedantic -Werror -pthread
4
5 all: $(OBJECTS)
6     gcc $(OBJECTS) -pthread -o mfind
7     gcc $(FLAGS) crazy_search.c -o crazy_search
8
9 mfind.o: mfind.c
10     gcc $(FLAGS) -c mfind.c
11
12 list.o: list.c list.h
13     gcc $(FLAGS) -c list.c
14
15 parser.o: parser.c parser.h list.h
16     gcc $(FLAGS) -c parser.c
17
18 debug: FLAGS+=-DDEBUG -g
19 debug: all
20
21 test: all
22     ./mfind -td -p2 . .. fail mfind
23
24 time: FLAGS+=-DTIME
25 time: all
26
27 memtest: all
28     valgrind ./mfind -td -p2 . .. mfind
29
30 clean:
31     rm -f mfind *.o
```