

# Operating systems – Assignment 3

## I/O Scheduling

Lennart Jern  
CS: ens16ljn

**Teacher**  
Ahmed Aley

January 7, 2017

# 1 Introduction

Disk access is significantly slower than any CPU operation and is often a bottle neck when it comes to performance. In this report I compare the performance of three I/O schedulers: “cfq”, “noop” and “deadline”. Additionally, all tests are done on two different types of disks: a USB flash drive and a hard disk drive (also connected via USB).

The benchmark used consists of a search for files in a hierarchy of directories. In other words, this report is focused on seek times since nothing is written to disk and just meta data (such as file name and directory content) is read.

## 2 Method

The benchmark consists of a program (`mfind`, see listing 4) that searches for files in a hierarchy of directories and prints out the time taken for each call to `readdir`. A script (`timer.sh`, see listing 1) is used to run `mfind` in four parallel processes, ten times for each of the three schedulers, and collects all timing data in log files. These logs are then processed by `stats.py` (see listing 2) in order to obtain some statistical properties of the data. The python library Pandas<sup>1</sup> proved very helpful in this regard.

All tests were run on my personal computer with the specifications seen in table 1. The drives used was a Kingston DataTraveler 1 GB and a Verbatim 500 GB portable 2.5” HDD. Both drives were connected to a USB 2.0 port.

## 3 Results

## 4 Final thoughts and lessons learned

It is quite clear that it is possible to spend a considerable amount of time just analyzing schedulers. The results are intriguing as they show clear differences in the behavior between the schedulers, and with several different tasks to compare, even more patterns would surely emerge.

---

<sup>1</sup><http://pandas.pydata.org/>

Component	Specification
OS:	Fedora 25
Kernel:	Linux 4.8.13-300.fc25.x86_64
CPU:	Intel Core i5-2500K CPU @ 3.7GHz
RAM:	7965MiB
GCC:	6.2.1
Bash:	4.3.43
Python:	3.5.2
Pandas:	0.18.1
Matplotlib:	1.5.3

Table 1: Test system specification.

	cfq	deadline	noop
count	$1.709 \cdot 10^5$	$1.709 \cdot 10^5$	$4.399 \cdot 10^{-4}$
mean	$4.546 \cdot 10^{-4}$	$4.478 \cdot 10^{-4}$	$1.395 \cdot 10^{-3}$
std	$1.465 \cdot 10^{-3}$	$1.396 \cdot 10^{-3}$	$2.100 \cdot 10^{-8}$
min	$2.600 \cdot 10^{-8}$	$2.600 \cdot 10^{-8}$	$8.400 \cdot 10^{-8}$
25%	$8.400 \cdot 10^{-8}$	$8.500 \cdot 10^{-8}$	$2.150 \cdot 10^{-7}$
50%	$2.160 \cdot 10^{-7}$	$2.140 \cdot 10^{-7}$	$1.845 \cdot 10^{-5}$
75%	$1.914 \cdot 10^{-5}$	$1.888 \cdot 10^{-5}$	$4.208 \cdot 10^{-2}$
max	$6.358 \cdot 10^{-2}$	$3.325 \cdot 10^{-2}$	$4.208 \cdot 10^{-2}$

Table 2: Kingston stats

A lesson learned is that one should think carefully about when and where to start and stop the timers. If measuring just the whole task, this is quite easy, but for the individual threads it gets more tricky. The thread that starts the other threads must also start the timers since the working thread may not get to start immediately. Similarly, the working threads must stop their own timers, since there might be a pause between the threads finishing and actually getting joined.

	cfq	deadline	noop
count	$1.709 \cdot 10^5$	$1.709 \cdot 10^5$	$4.413 \cdot 10^{-4}$
mean	$2.892 \cdot 10^{-4}$	$4.626 \cdot 10^{-4}$	$3.030 \cdot 10^{-3}$
std	$2.018 \cdot 10^{-3}$	$3.082 \cdot 10^{-3}$	$2.400 \cdot 10^{-8}$
min	$2.300 \cdot 10^{-8}$	$2.600 \cdot 10^{-8}$	$1.230 \cdot 10^{-7}$
25%	$1.170 \cdot 10^{-7}$	$1.210 \cdot 10^{-7}$	$2.590 \cdot 10^{-7}$
50%	$2.490 \cdot 10^{-7}$	$2.610 \cdot 10^{-7}$	$2.366 \cdot 10^{-6}$
75%	$2.036 \cdot 10^{-6}$	$2.320 \cdot 10^{-6}$	$8.346 \cdot 10^{-2}$
max	$1.259 \cdot 10^{-1}$	$8.426 \cdot 10^{-2}$	$8.346 \cdot 10^{-2}$

Table 3: Verbatim stats

## A Code listings

Listing 1: timer.sh

```
1  #!/bin/bash
2
3  # timer.sh
4  #
5  # A timer script to measure the differences between i/o schedulers
6  #
7  # Author: Lennart Jern (ens16ljn@cs.umu.se)
8
9  # Clean up old results
10 rm ../data/*.log
11
12 SCHEDULERS="cfq_noop_deadline"
13 DEVICE="/sys/block/sdc/queue/scheduler"
14 # Starting directory and expression to search for
15 # START="/run/media/lennart/KINGSTON/test_files_expression"
16 START="/run/media/lennart/Verbatim/test_files_expression"
17 # MNT="/run/media/lennart/KINGSTON"
18 MNT="/run/media/lennart/Verbatim"
19
20 for S in $SCHEDULERS
21 do
22     echo $S | sudo tee $DEVICE
23     echo "Scheduler: `cat $DEVICE`"
24     # LINE=""
25     # Time the commands 10 times
26     for i in $(seq 1 10)
27     do
28         # Unmount and mount to clear all cache
29         sync
30         sleep 1
31         sudo umount -f $MNT
32         sudo rm -d $MNT
33         sudo mkdir $MNT
34         sudo mount /dev/sdc1 $MNT
35
36         # We use 4 parallel commands that store their respective times in
37         # separate log files
38         COMMAND1="./mfind $START >> ../data/$S-1.log"
39         COMMAND2="./mfind $START >> ../data/$S-2.log"
40         COMMAND3="./mfind $START >> ../data/$S-3.log"
41         COMMAND4="./mfind $START >> ../data/$S-4.log"
42
43         eval $COMMAND1 &
44         eval $COMMAND2 &
45         eval $COMMAND3 &
46         eval $COMMAND4 &
47         # Wait for all commands to finish
48         wait
49
50         # A little progress report
51         echo "Run $i done."
52     done
53 done
54
55 # Restore cfq scheduler
56 echo "cfq" | sudo tee $DEVICE
```

---

## Listing 2: stats.py

```
1  """
2  stats.py
3
4  Process the data produced by timer.sh by calculating the
5  medians, max values and min values for each scheduler.
6  Also plots the density curves.
7
8  Author: Lennart Jern (ens16ljn@cs.umu.se)
9  """
10
11 import pandas as pd
12 import re
13 import matplotlib.pyplot as plt
14
15 def produce_stats():
16     """
17     Read the data from files, calculate statistical values and make a plot.
18     """
19     base = "../data/"
20     ext = ".csv"
21     header=("cfq", "deadline", "noop")
22
23     # Get individual read times as a DataFrame
24     df = get_read_times()
25
26     stats = df.describe()
27     # Escape per cent chars
28     idx = ['count', 'mean', 'std', 'min', '25%', '50%', '75%', 'max']
29     stats = stats.set_index([idx])
30
31     # Save stats as csv
32     stats.to_csv(base+"stats.csv", header=header, float_format="%.3e")
33
34     # Plot and save the density curves
35     ax = df.plot.kde()
36     ax.set_xlabel("Time_(s)")
37     ax.set_xlim([0, 0.006])
38     fig = ax.get_figure()
39     fig.savefig(base+"density.pdf")
40
41
42 def collect_read_times(file_name):
43     """Read thread times from a file."""
44     f = open(file_name)
45     times = []
46     # Regular expression to find floats
47     time = re.compile("(\\d+\\.\\d+)")
48
49     for line in f:
50         match = time.match(line)
51
52         if (match):
53             t = float(match.group(1))
54             times.append(t)
55
56     return times
57
```

```

58 def get_read_times():
59     """Collect timing information about all schedulers in a DataFrame."""
60     schedulers = ["cfq", "deadline", "noop"]
61     base = "../data/"
62     ext = ".log"
63     header=("cfq", "deadline", "noop")
64
65     # Collect all times in one file
66     times = {key: [] for key in schedulers}
67     for s in schedulers:
68         # We have 4 parallel log files for each scheduler
69         for i in [1,2,3,4]:
70             f = base+s+"-"+str(i)+ext
71             times[s].extend(collect_read_times(f))
72     # Return a DataFrame with all timing data
73     df = pd.DataFrame(times)
74     return df
75
76 # Collect statistical data
77 produce_stats()

```

Listing 3: generate\_test\_files.sh

```

1  #!/bin/bash
2
3  # File: generate_test_files.sh
4  # Author: Lennart Jern - ens16ljn@cs.umu.se
5  #
6  # Generate a file tree to do tests on.
7
8  # Starting directory
9  # START=/media/removable/KINGSTON
10 START=/run/media/lennart/Verbatim
11
12 # Move to correct directory/device
13 cd $START
14
15 # Remove tree if existent
16 rm -r test_files
17
18 # Create directory to hold all test files
19 mkdir test_files
20 cd test_files
21
22 # Split up the files between a few directories
23 for DIR in a b c d e; do
24     mkdir $DIR
25     cd $DIR
26     # Create empty files
27     for F in $(seq 1 20); do
28         touch $F
29     done
30
31     for D in f g h i j; do
32         mkdir $D
33         cd $D
34     done
35
36     cd "$START/test_files"
37 done
38

```

```

39 # Big files
40 mkdir bigs
41 cd bigs
42
43 # Generate files with lots of zeros...
44 for F in $(seq 1 5); do
45     head -c 50M < /dev/zero > "file$F"
46 done
47
48 cd "$START/test_files"
49 # Generate deep folders
50 for i in $(seq 1 10); do
51     mkdir "deep$i"
52     cd "deep$i"
53     for D in $(seq 1 100); do
54         mkdir "dir$D"
55         cd "dir$D"
56     done
57     cd "$START/test_files"
58 done

```

Listing 4: mfind.c

```

1 /**
2  * File: mfind.c
3  * Author: Lennart Jern - ens16ljn@cs.umu.se
4  *
5  * Usage: ./mfind [-t {d|f|l}] [-p nrthr] start1 [start2 ...] name
6  *
7  * mfind can search after files, links and directories from given start paths.
8  * The search can be done with more than one thread by specifying the flag
9  * '-p#', where # is the number of threads to use.
10  *
11  */
12 #define _GNU_SOURCE
13
14 #include <stdio.h>
15 #include <stdlib.h>
16 #include <errno.h>
17 #include <string.h>
18 #include <dirent.h>
19 #include <sys/stat.h>
20 #include <time.h>           // timing
21 #include "parser.h"        // Includes list.h
22
23 #define ONE_OVER_BILLION 1E-9
24
25 void *find_file(void *s_data);
26 int search_path(SearchData *data, char *path);
27 int search_directory(SearchData *search_data, DIR *dir, char *path);
28 int get_dirent(struct dirent *priv_dirent, DIR *dir);
29 void process_file(char *file_path, char *name,
30                  struct stat f_stat, SearchData *data);
31 int add_dir(LinkedList *list, char *dir_path);
32 void check_starting_dirs(SearchData *search_data);
33 void print_path(void *path);
34 void delete_path(void *path);
35 void SearchData_delete(SearchData *s_data);
36
37 /**
38  * main - parse arguments, do the search and then clean up.

```



```

39  * @param argc -- number of arguments
40  * @param argv -- array of arguments
41  * @return      0 if everything went well, a positive int otherwise
42  */
43  int main(int argc, char *argv[]) {
44      int ret = 0;
45
46      SearchData *search_data = parse_arguments(argc, argv);
47      check_starting_dirs(search_data);
48
49      #ifdef DEBUG
50          printf("#_Search_data\n=====\\n");
51          printf("#_Threads:_%d\\n", search_data->num_threads);
52          printf("#_Type:_%c\\n", search_data->type);
53          printf("#_Needle:_%s\\n", search_data->needle);
54          List_print(search_data->directories, print_path);
55          printf("=====\\n\\n");
56      #endif
57
58      search_data->num_searchers = 0;
59
60      find_file(search_data);
61
62      // Check for errors
63      ret = search_data->error;
64      // Free allocated memory
65      SearchData_delete(search_data);
66      return ret;
67  }
68
69  /**
70   * find_file - search for files and directories
71   * @param s_data -- search data, containing needle to look for and list of
72   *                directories to look in
73   */
74  void *find_file(void *search_data) {
75      unsigned int reads = 0;
76      SearchData *data = search_data;
77      char *path = NULL;
78      int error = 0;
79
80      // Keep searching while there are dirs in the list.
81      while((path = (char *)List_get(data->directories)) != NULL) {
82
83          data->num_searchers++;
84
85          reads++;
86          if (search_path(data, path) != 0) {
87              perror(path);
88              // We don't consider permission denied or missing dir as errors
89              // error = 1;
90          }
91
92          delete_path(path);
93          data->num_searchers--;
94      } // End while. No more dirs to search and all threads done.
95      // Make sure caller knows if there were errors.
96      data->error = error;
97      printf("Reads:_%d\\n", reads);
98      return NULL;
99  }
100

```

```

101 /**
102  * search_path - open and search the directory given by path
103  * @param data -- SearchData (what to search for)
104  * @param path -- path to directory to search
105  * @return      0 on successful search, -1 if there were errors
106  */
107 int search_path(SearchData *data, char *path) {
108     // Open the directory. If it fails, clean up and continue with the next one.
109     DIR *dir = opendir(path);
110     int ret = 0;
111
112     if (dir == NULL) {
113         ret = -1;
114         return ret;
115     }
116
117     // Check for matches in the dir
118     if (search_directory(data, dir, path) != 0) {
119         ret = -1;
120     }
121
122     if (closedir(dir) != 0) {
123         perror("closedir");
124     }
125     return ret;
126 }
127
128 /**
129  * search_directory - check all files and folders in dir for matches and
130  * add folders to the list.
131  * @param search_data -- data regarding the search
132  * @param dir          -- dir to look in
133  * @param path         -- path to the dir (used for printing)
134  * @return             0 if everything went well, a poitive int otherwise.
135  */
136 int search_directory(SearchData *search_data, DIR *dir, char *path) {
137     struct dirent *priv_dirent;
138     struct stat f_stat;
139     char *file_path = NULL;
140     int at_end = 0;
141     int ret = 0;
142
143     while (at_end != 1) {
144         priv_dirent = malloc(sizeof(struct dirent));
145         if (priv_dirent == NULL) {
146             perror("malloc");
147             exit(EXIT_FAILURE);
148         }
149
150         at_end = get_dirent(priv_dirent, dir);
151         if (at_end != 0) {
152             // Either error or end of dir
153             if (at_end == -1) {
154                 ret++;
155             }
156             free(priv_dirent);
157             continue;
158         }
159
160         // Build file path string
161         if (asprintf(&file_path, "%s/%s", path, priv_dirent->d_name) == -1) {
162             fprintf(stderr, "Error: asprintf failed. Unable to set file_path.\n");

```

```

163         free(priv_dirent);
164         ret++;
165         continue;
166     }
167     // Get stats (file type)
168     if (lstat(file_path, &f_stat) != 0) {
169         perror(file_path);
170         free(priv_dirent);
171         free(file_path);
172         ret++;
173         continue;
174     }
175
176     // Print matches.
177     process_file(file_path, priv_dirent->d_name, f_stat, search_data);
178
179     // Add directories to the list (not . and ..)
180     if (S_ISDIR(f_stat.st_mode)
181         && strcmp(priv_dirent->d_name, ".") != 0
182         && strcmp(priv_dirent->d_name, "..") != 0) {
183
184         if (add_dir(search_data->directories, file_path) != 1) {
185             fprintf(stderr, "Failed to add directory to list.\n");
186             return -1;
187         }
188     }
189
190     free(file_path);
191     file_path = NULL;
192     free(priv_dirent);
193     priv_dirent = NULL;
194 }
195 return ret;
196 }
197
198 /**
199  * get_dirent - copy the next dirent in dir to priv_dirent in a thread safe way.
200  * This private dirent is safe to use in a multi thread environment.
201  * @param priv_dirent -- pointer to dirent where the dirent will be saved.
202  * @param dir          -- dir to read from
203  * @return             -1 on error, 1 when the last element was read and
204  *                     0 otherwise
205  */
206 int get_dirent(struct dirent *priv_dirent, DIR *dir) {
207     struct dirent *dirent;
208     errno = 0;
209
210     // Starting time
211     struct timespec start;
212     // Time when finished
213     struct timespec end;
214     clock_gettime(CLOCK_REALTIME, &start);
215
216     dirent = readdir(dir);
217
218     // Get the time when finished
219     clock_gettime(CLOCK_REALTIME, &end);
220     // Calculate time it took
221     double time_taken = (end.tv_sec - start.tv_sec)
222         + (end.tv_nsec - start.tv_nsec)
223         * ONE_OVER_BILLION;
224     printf("%.12lf\n", time_taken);

```

```

225
226     if (errno != 0) {
227         perror("readdir");
228         return -1;
229     } else if (dirent == NULL) {
230         // No more files to read
231         return 1;
232     }
233     // Copy dirent to private memory
234     memcpy(priv_dirent, dirent, sizeof(struct dirent));
235     return 0;
236 }
237
238 /**
239  * process_file - print out matching file.
240  * @param file_path -- the path to the file
241  * @param name      -- name of the file
242  * @param f_stat    -- file stats
243  * @param data      -- SearchData (what type and name are we looking for?)
244  */
245 void process_file(char *file_path, char *name,
246                  struct stat f_stat, SearchData *data) {
247     int match = 0;
248     char type = data->type;
249     // Is the name matching?
250     if (strcmp(name, data->needle) == 0) {
251         match = 1;
252     }
253
254     // Check type, print if we have a match
255     if (S_ISDIR(f_stat.st_mode)) {
256         if (match == 1 && (type == 'd' || type == '\0')) {
257             printf("%s\n", file_path);
258         }
259     } else if (S_ISREG(f_stat.st_mode)) {
260         if (match == 1 && (type == 'f' || type == '\0')) {
261             printf("%s\n", file_path);
262         }
263     } else if (S_ISLNK(f_stat.st_mode)) {
264         if (match == 1 && (type == 'l' || type == '\0')) {
265             printf("%s\n", file_path);
266         }
267     }
268 }
269
270 /**
271  * add_dir - add a directory to the list in a thread safe manner
272  * @param list    -- list to add to
273  * @param dir_path -- path to the directory
274  * @return        1 if the dir was added, 0 if addition failed.
275  */
276 int add_dir(LinkedList *list, char *dir_path) {
277     char *new_dir = malloc(strlen(dir_path)+1);
278     if (new_dir == NULL) {
279         perror("malloc");
280         exit(EXIT_FAILURE);
281     }
282     strcpy(new_dir, dir_path);
283
284     if (List_append(list, (void *)new_dir) != 1) {
285         return 0;
286     }

```

```

287     return 1;
288 }
289
290 /**
291  * check_starting_dirs - check if the starting dirs match the search criterias
292  * @param search_data -- data egrading the search
293  */
294 void check_starting_dirs(SearchData *search_data) {
295     char *path;
296     struct stat f_stat;
297     LinkedList *checked_dirs = List_init();
298
299     // Check all starting dirs for matches
300     while((path = (char *)List_get(search_data->directories)) != NULL) {
301         if (lstat(path, &f_stat) != 0) {
302             perror(path);
303             continue;
304         }
305         // Print if there is a match
306         process_file(path, basename(path), f_stat, search_data);
307
308         // Add the checked dir to the new list
309         char *new_dir = malloc(strlen(path)+1);
310         if (new_dir == NULL) {
311             perror("malloc");
312             exit(EXIT_FAILURE);
313         }
314         strcpy(new_dir, path);
315
316         if (List_append(checked_dirs, (void *)new_dir) == 0) {
317             fprintf(stderr, "Could not add path to list.\n");
318             search_data->error++;
319         }
320         free(path);
321     }
322     // Delete the old list
323     List_delete(search_data->directories, delete_path);
324     // Add the checked dirs
325     search_data->directories = checked_dirs;
326 }
327
328 /**
329  * print_path - print out a path
330  * @param path -- a void pointer to a path string
331  */
332 void print_path(void *path) {
333     char *str = (char *)path;
334     printf("%s\n", str);
335 }
336
337 /**
338  * delete_path - delete and free any memory occupied by path
339  * @param path -- path to be freed
340  */
341 void delete_path(void *path) {
342     free(path);
343 }
344
345 /**
346  * SearchData_delete - free all memory allocated for s_data
347  * @param s_data -- SearchData to free
348  */

```

```

349 void SearchData_delete(SearchData *s_data) {
350     free(s_data->needle);
351     List_delete(s_data->directories, delete_path);
352     free(s_data);
353 }

```

### Listing 5: Makefile

```

1  SOURCE=mfind.c list.c
2  OBJECTS=mfind.o list.o parser.o
3  FLAGS=-std=c11 -Wall -pedantic -Werror -pthread
4
5  all: $(OBJECTS)
6      gcc $(OBJECTS) -pthread -o mfind
7      gcc $(FLAGS) crazy_search.c -o crazy_search
8
9  mfind.o: mfind.c
10     gcc $(FLAGS) -c mfind.c
11
12  list.o: list.c list.h
13     gcc $(FLAGS) -c list.c
14
15  parser.o: parser.c parser.h list.h
16     gcc $(FLAGS) -c parser.c
17
18  debug: FLAGS+=-DDEBUG -g
19  debug: all
20
21  test: all
22     ./mfind -td -p2 . .. fail mfind
23
24  time: FLAGS+=-DTIME
25  time: all
26
27  memtest: all
28     valgrind ./mfind -td -p2 . .. mfind
29
30  clean:
31     rm -f mfind *.o

```