

# Operating systems – Assignment 3

## I/O Scheduling

Lennart Jern  
CS: ens16ljn

**Teacher**  
Ahmed Aley

January 8, 2017

# 1 Introduction

Disk access is significantly slower than any CPU operation and is often a bottle neck when it comes to performance. In this report I compare the performance of three I/O schedulers: “cfq”, “noop” and “deadline”. Additionally, all tests are done on two different types of disks: a USB flash drive and a hard disk drive (also connected via USB).

The benchmark used consists of a search for files in a hierarchy of directories. In other words, this report is focused on seek times since nothing is written to disk and just meta data (such as file name and directory content) is read.

## 2 Method

The benchmark consists of a program (`mfind`, see listing 1) that searches for files in a hierarchy of directories and prints out the time taken for each call to `readdir`. A script (`timer.sh`, see listing 2) is used to run `mfind` in four parallel processes, ten times for each of the three schedulers, and collects all timing data in log files. The partition is unmounted between each run to make sure that no files are cached.

Processing of the data is done by `stats.py` (see listing 3) in order to obtain some statistical properties. The python library Pandas<sup>1</sup> proved very helpful in this regard.

**Remark 1.** *The program `mfind` was originally made for another assignment and was here simplified and modified to measure the time required for each I/O request. In addition to `mfind.c` the program also consists of a parser (`parser.c` and `parser.h` in listings 4 and 5) and the linked list that keeps track of what directories to search (`list.c` and `list.h` in listings 6 and 7). The complete program can be compiled using the make file in listing 8.*

A script was also used to create the tree of searched directories in order to make the benchmark easier to reproduce. This script can be found in listing 9.

All tests were run on my personal computer with the specifications seen in table 1. The drives used was a Kingston DataTraveler 1 GB and a Verbatim 500 GB portable 2.5” HDD. Both drives were connected to a USB 2.0

---

<sup>1</sup><http://pandas.pydata.org/>

port.

Component	Specification
OS:	Fedora 25
Kernel:	Linux 4.8.13-300.fc25.x86_64
CPU:	Intel Core i5-2500K CPU @ 3.7GHz
RAM:	7965MiB
GCC:	6.2.1
Bash:	4.3.43
Python:	3.5.2
Pandas:	0.18.1
Matplotlib:	1.5.3

Table 1: Test system specification.

### 3 Results

The calculated statistical properties of the timing data can be seen in tables 2 and 3. Only slight differences can be seen between the schedulers for the flash drive. Most noticeably, the maximum time required is considerably greater for cfq than for the other schedulers. This also reflects on a greater standard deviation.

	cfq	deadline	noop
count	$1.709 \cdot 10^5$	$1.709 \cdot 10^5$	$1.709 \cdot 10^5$
mean	$4.546 \cdot 10^{-4}$	$4.478 \cdot 10^{-4}$	$4.399 \cdot 10^{-4}$
std	$1.465 \cdot 10^{-3}$	$1.396 \cdot 10^{-3}$	$1.395 \cdot 10^{-3}$
min	$2.600 \cdot 10^{-8}$	$2.600 \cdot 10^{-8}$	$2.100 \cdot 10^{-8}$
25%	$8.400 \cdot 10^{-8}$	$8.500 \cdot 10^{-8}$	$8.400 \cdot 10^{-8}$
50%	$2.160 \cdot 10^{-7}$	$2.140 \cdot 10^{-7}$	$2.150 \cdot 10^{-7}$
75%	$1.914 \cdot 10^{-5}$	$1.888 \cdot 10^{-5}$	$1.845 \cdot 10^{-5}$
max	$6.358 \cdot 10^{-2}$	$3.325 \cdot 10^{-2}$	$4.208 \cdot 10^{-2}$

Table 2: Statistical properties for the Kingston flash drive.  
The time values are measured in seconds.

Looking at the means and medians (50 % quantile), we see that noop has the lowest mean, but is beaten by deadline when it comes to the median. Noop also has the lowest minimum time required.

The picture is quite different when looking at table 3. This is for a traditional spinning hard drive, so the cfq scheduler get to shine. The cfq scheduler has the best mean, minimum, median and 75 % quantile. All this comes at a cost, of course, the maximum time is far greater than for the other schedulers. Despite this, however, the cfq still manages to have the lowest standard deviation among the three.

	cfq	deadline	noop
count	$1.709 \cdot 10^5$	$1.709 \cdot 10^5$	$1.709 \cdot 10^5$
mean	$2.892 \cdot 10^{-4}$	$4.626 \cdot 10^{-4}$	$4.413 \cdot 10^{-4}$
std	$2.018 \cdot 10^{-3}$	$3.082 \cdot 10^{-3}$	$3.030 \cdot 10^{-3}$
min	$2.300 \cdot 10^{-8}$	$2.600 \cdot 10^{-8}$	$2.400 \cdot 10^{-8}$
25%	$1.170 \cdot 10^{-7}$	$1.210 \cdot 10^{-7}$	$1.230 \cdot 10^{-7}$
50%	$2.490 \cdot 10^{-7}$	$2.610 \cdot 10^{-7}$	$2.590 \cdot 10^{-7}$
75%	$2.036 \cdot 10^{-6}$	$2.320 \cdot 10^{-6}$	$2.366 \cdot 10^{-6}$
max	$1.259 \cdot 10^{-1}$	$8.426 \cdot 10^{-2}$	$8.346 \cdot 10^{-2}$

Table 3: Statistical properties for the Verbatim HDD. The time values are measured in seconds.

A comparison of the two drives reveals that the scheduler should be chosen depending on the type of drive for best performance. On a traditional spinning drive, there is more to win by ordering the requests in an efficient way. However, the ordering does not matter much on a solid state drive, since there is no physical movement required. In that case it is better not to spend time finding the “perfect order” and thus delay some requests.

In order to improve the noop scheduler on traditional HDDs we would have to allow a higher maximum time. This would give the scheduler a better possibility to control the ordering of the requests.

Similarly, the cfq scheduler would probably be faster on flash drives if it just sent all requests in the order they came in, instead of reordering them. The deadline scheduler is similar in performance to the noop scheduler on both drives and could probably improve its performance on HDDs in the same way. My guess is that the deadline and noop schedulers would show greater differences if the requests were longer.

## 4 Final thoughts and lessons learned

Caching of data gave some problems at first, but this could easily be solve by unmounting the drive between each run. However, the unmounting was quite problematic for the HDD since it would sometimes remain busy even after a `sync` call. To solve this, a one second sleep was added to the script.

The problem with `sync` and `umount` made me quite suspicious of the data gathered. I am uncertain about what really should be timed and how to make sure that the I/O requests are actually done on the disk and not just in some cache. Perhaps more fine tuned control would be needed to solve these problems, with calls to `syncfs` periodically. But that makes it hard to determine the actual time taken for individual requests.

One thing is clear: the I/O scheduler can definitely have a great impact on performance and it can be quite tricky to accurately measure the performance of these schedulers. Some theoretical background would have been interesting and helpful if time would have allowed for that during the lectures.

## A Code listings

Listing 1: mfind.c

```
1  /**
2   * File: mfind.c
3   * Author: Lennart Jern - ens16ljn@cs.umu.se
4   *
5   * Usage: ./mfind [-t {d|f|l}] [-p nrthr] start1 [start2 ...] name
6   *
7   * mfind can search after files, links and directories from given start paths.
8   * The search can be done with more than one thread by specifying the flag
9   * '-p#', where # is the number of threads to use.
10  *
11  */
12  #define _GNU_SOURCE
13
14  #include <stdio.h>
15  #include <stdlib.h>
16  #include <errno.h>
17  #include <string.h>
18  #include <dirent.h>
19  #include <sys/stat.h>
20  #include <time.h>           // timing
21  #include "parser.h"        // Includes list.h
22
23  #define ONE_OVER_BILLION 1E-9
24
25  void *find_file(void *s_data);
26  int search_path(SearchData *data, char *path);
27  int search_directory(SearchData *search_data, DIR *dir, char *path);
28  int get_dirent(struct dirent *priv_dirent, DIR *dir);
29  void process_file(char *file_path, char *name,
30                  struct stat f_stat, SearchData *data);
31  int add_dir(LinkedList *list, char *dir_path);
32  void check_starting_dirs(SearchData *search_data);
33  void print_path(void *path);
34  void delete_path(void *path);
35  void SearchData_delete(SearchData *s_data);
36
37  /**
38   * main - parse arguments, do the search and then clean up.
39   * @param argc -- number of arguments
40   * @param argv -- array of arguments
41   * @return      0 if everything went well, a positive int otherwise
42   */
43  int main(int argc, char *argv[]) {
44      int ret = 0;
45
46      SearchData *search_data = parse_arguments(argc, argv);
47      check_starting_dirs(search_data);
48
49      #ifdef DEBUG
50          printf("#_Search_data\n=====\\n");
51          printf("#_Threads:_%d\\n", search_data->num_threads);
52          printf("#_Type:_%c\\n", search_data->type);
53          printf("#_Needle:_%s\\n", search_data->needle);
54          List_print(search_data->directories, print_path);
55          printf("=====\\n\\n");
56      #endif
57  }
```

```

58     search_data->num_searchers = 0;
59
60     find_file(search_data);
61
62     // Check for errors
63     ret = search_data->error;
64     // Free allocated memory
65     searchData_delete(search_data);
66     return ret;
67 }
68
69 /**
70  * find_file - search for files and directories
71  * @param s_data -- search data, containing needle to look for and list of
72  *                directories to look in
73  */
74 void *find_file(void *search_data) {
75     unsigned int reads = 0;
76     searchData *data = search_data;
77     char *path = NULL;
78     int error = 0;
79
80     // Keep searching while there are dirs in the list.
81     while((path = (char *)List_get(data->directories)) != NULL) {
82
83         data->num_searchers++;
84
85         reads++;
86         if (search_path(data, path) != 0) {
87             perror(path);
88             // We don't consider permission denied or missing dir as errors
89             // error = 1;
90         }
91
92         delete_path(path);
93         data->num_searchers--;
94     } // End while. No more dirs to search and all threads done.
95     // Make sure caller knows if there were errors.
96     data->error = error;
97     printf("Reads: %d\n", reads);
98     return NULL;
99 }
100
101 /**
102  * search_path - open and search the directory given by path
103  * @param data -- searchData (what to search for)
104  * @param path -- path to directory to search
105  * @return      0 on successful search, -1 if there were errors
106  */
107 int search_path(SearchData *data, char *path) {
108     // Open the directory. If it fails, clean up and continue with the next one.
109     DIR *dir = opendir(path);
110     int ret = 0;
111
112     if (dir == NULL) {
113         ret = -1;
114         return ret;
115     }
116
117     // Check for matches in the dir
118     if (search_directory(data, dir, path) != 0) {
119         ret = -1;

```

```

120     }
121
122     if (closedir(dir) != 0) {
123         perror("closedir");
124     }
125     return ret;
126 }
127
128 /**
129  * search_directory - check all files and folders in dir for matches and
130  * add folders to the list.
131  * @param search_data -- data regarding the search
132  * @param dir          -- dir to look in
133  * @param path         -- path to the dir (used for printing)
134  * @return             0 if everything went well, a poitive int otherwise.
135  */
136 int search_directory(SearchData *search_data, DIR *dir, char *path) {
137     struct dirent *priv_dirent;
138     struct stat f_stat;
139     char *file_path = NULL;
140     int at_end = 0;
141     int ret = 0;
142
143     while (at_end != 1) {
144         priv_dirent = malloc(sizeof(struct dirent));
145         if (priv_dirent == NULL) {
146             perror("malloc");
147             exit(EXIT_FAILURE);
148         }
149
150         at_end = get_dirent(priv_dirent, dir);
151         if (at_end != 0) {
152             // Either error or end of dir
153             if (at_end == -1) {
154                 ret++;
155             }
156             free(priv_dirent);
157             continue;
158         }
159
160         // Build file path string
161         if (asprintf(&file_path, "%s/%s", path, priv_dirent->d_name) == -1) {
162             fprintf(stderr, "Error: asprintf failed. Unable to set file_path.\n");
163             free(priv_dirent);
164             ret++;
165             continue;
166         }
167         // Get stats (file type)
168         if (lstat(file_path, &f_stat) != 0) {
169             perror(file_path);
170             free(priv_dirent);
171             free(file_path);
172             ret++;
173             continue;
174         }
175
176         // Print matches.
177         process_file(file_path, priv_dirent->d_name, f_stat, search_data);
178
179         // Add directories to the list (not . and ..)
180         if (S_ISDIR(f_stat.st_mode)
181             && strcmp(priv_dirent->d_name, ".") != 0

```



```

182         && strcmp(priv_dirent->d_name, "..") != 0) {
183
184             if (add_dir(search_data->directories, file_path) != 1) {
185                 fprintf(stderr, "Failed to add directory to list.\n");
186                 return -1;
187             }
188         }
189
190         free(file_path);
191         file_path = NULL;
192         free(priv_dirent);
193         priv_dirent = NULL;
194     }
195     return ret;
196 }
197
198 /**
199  * get_dirent - copy the next dirent in dir to priv_dirent in a thread safe way.
200  * This private dirent is safe to use in a multi thread environment.
201  * @param priv_dirent -- pointer to dirent where the dirent will be saved.
202  * @param dir         -- dir to read from
203  * @return            -1 on error, 1 when the last element was read and
204  *                    0 otherwise
205  */
206 int get_dirent(struct dirent *priv_dirent, DIR *dir) {
207     struct dirent *dirent;
208     errno = 0;
209
210     // Starting time
211     struct timespec start;
212     // Time when finished
213     struct timespec end;
214     clock_gettime(CLOCK_REALTIME, &start);
215
216     dirent = readdir(dir);
217
218     // Get the time when finished
219     clock_gettime(CLOCK_REALTIME, &end);
220     // Calculate time it took
221     double time_taken = (end.tv_sec - start.tv_sec)
222         + (end.tv_nsec - start.tv_nsec)
223         * ONE_OVER_BILLION;
224     printf("%.12lf\n", time_taken);
225
226     if (errno != 0) {
227         perror("readdir");
228         return -1;
229     } else if (dirent == NULL) {
230         // No more files to read
231         return 1;
232     }
233     // Copy dirent to private memory
234     memcpy(priv_dirent, dirent, sizeof(struct dirent));
235     return 0;
236 }
237
238 /**
239  * process_file - print out matching file.
240  * @param file_path -- the path to the file
241  * @param name       -- name of the file
242  * @param f_stat      -- file stats
243  * @param data        -- SearchData (what type and name are we looking for?)

```

```

244  */
245  void process_file(char *file_path, char *name,
246                  struct stat f_stat, SearchData *data) {
247      int match = 0;
248      char type = data->type;
249      // Is the name matching?
250      if (strcmp(name, data->needle) == 0) {
251          match = 1;
252      }
253
254      // Check type, print if we have a match
255      if (S_ISDIR(f_stat.st_mode)) {
256          if (match == 1 && (type == 'd' || type == '\0')) {
257              printf("%s\n", file_path);
258          }
259      } else if (S_ISREG(f_stat.st_mode)) {
260          if (match == 1 && (type == 'f' || type == '\0')) {
261              printf("%s\n", file_path);
262          }
263      } else if (S_ISLNK(f_stat.st_mode)) {
264          if (match == 1 && (type == 'l' || type == '\0')) {
265              printf("%s\n", file_path);
266          }
267      }
268  }
269
270  /**
271   * add_dir - add a directory to the list in a thread safe manner
272   * @param list    -- list to add to
273   * @param dir_path -- path to the directory
274   * @return        1 if the dir was added, 0 if addition failed.
275   */
276  int add_dir(LinkedList *list, char *dir_path) {
277      char *new_dir = malloc(strlen(dir_path)+1);
278      if (new_dir == NULL) {
279          perror("malloc");
280          exit(EXIT_FAILURE);
281      }
282      strcpy(new_dir, dir_path);
283
284      if (List_append(list, (void *)new_dir) != 1) {
285          return 0;
286      }
287      return 1;
288  }
289
290  /**
291   * check_starting_dirs - check if the starting dirs match the search criterias
292   * @param search_data -- data egarding the search
293   */
294  void check_starting_dirs(SearchData *search_data) {
295      char *path;
296      struct stat f_stat;
297      LinkedList *checked_dirs = List_init();
298
299      // Check all starting dirs for matches
300      while((path = (char *)List_get(search_data->directories)) != NULL) {
301          if (lstat(path, &f_stat) != 0) {
302              perror(path);
303              continue;
304          }
305          // Print if there is a match

```

```

306     process_file(path, basename(path), f_stat, search_data);
307
308     // Add the checked dir to the new list
309     char *new_dir = malloc(strlen(path)+1);
310     if (new_dir == NULL) {
311         perror("malloc");
312         exit(EXIT_FAILURE);
313     }
314     strcpy(new_dir, path);
315
316     if (List_append(checkered_dirs, (void *)new_dir) == 0) {
317         fprintf(stderr, "Could not add path to list.\n");
318         search_data->error++;
319     }
320     free(path);
321 }
322 // Delete the old list
323 List_delete(search_data->directories, delete_path);
324 // Add the checked dirs
325 search_data->directories = checked_dirs;
326 }
327
328 /**
329  * print_path - print out a path
330  * @param path -- a void pointer to a path string
331  */
332 void print_path(void *path) {
333     char *str = (char *)path;
334     printf("%s\n", str);
335 }
336
337 /**
338  * delete_path - delete and free any memory occupied by path
339  * @param path -- path to be freed
340  */
341 void delete_path(void *path) {
342     free(path);
343 }
344
345 /**
346  * searchData_delete - free all memory allocated for s_data
347  * @param s_data -- searchData to free
348  */
349 void searchData_delete(SearchData *s_data) {
350     free(s_data->needle);
351     List_delete(s_data->directories, delete_path);
352     free(s_data);
353 }

```

Listing 2: timer.sh

```

1  #!/bin/bash
2
3  # timer.sh
4  #
5  # A timer script to measure the differences between i/o schedulers
6  #
7  # Author: Lennart Jern (ens16ljn@cs.umu.se)
8
9  # Clean up old results
10 rm ../data/*.log

```

```

11
12 SCHEDULERS="cfq_noop_deadline"
13 DEVICE="/sys/block/sdc/queue/scheduler"
14 # Starting directory and expression to search for
15 # START="/run/media/lennart/KINGSTON/test_files expression"
16 START="/run/media/lennart/Verbatim/test_files expression"
17 # MNT="/run/media/lennart/KINGSTON"
18 MNT="/run/media/lennart/Verbatim"
19
20 for S in $SCHEDULERS
21 do
22     echo $S | sudo tee $DEVICE
23     echo "Scheduler: `cat $DEVICE`"
24     # LINE=""
25     # Time the commands 10 times
26     for i in $(seq 1 10)
27     do
28         # Unmount and mount to clear all cache
29         sync
30         sleep 1
31         sudo umount -f $MNT
32         sudo rm -d $MNT
33         sudo mkdir $MNT
34         sudo mount /dev/sdc1 $MNT
35
36         # We use 4 parallel commands that store their respective times in
37         # separate log files
38         COMMAND1="./mfind $START >> ../data/$S-1.log"
39         COMMAND2="./mfind $START >> ../data/$S-2.log"
40         COMMAND3="./mfind $START >> ../data/$S-3.log"
41         COMMAND4="./mfind $START >> ../data/$S-4.log"
42
43         eval $COMMAND1 &
44         eval $COMMAND2 &
45         eval $COMMAND3 &
46         eval $COMMAND4 &
47         # Wait for all commands to finish
48         wait
49
50         # A little progress report
51         echo "Run $i done."
52     done
53 done
54
55 # Restore cfq scheduler
56 echo "cfq" | sudo tee $DEVICE
57

```

Listing 3: stats.py

```

1 """
2 stats.py
3
4 Process the data produced by timer.sh by calculating the
5 medians, max values and min values for each scheduler.
6
7 Author: Lennart Jern (ens16ljn@cs.umu.se)
8 """
9
10 import pandas as pd
11 import re

```

```

12
13 def produce_stats():
14     """
15     Read the data from files, calculate statistical values and make a plot.
16     """
17     base = "../data/"
18     ext = ".csv"
19     header=("cfq", "deadline", "noop")
20
21     # Get individual read times as a DataFrame
22     df = get_read_times()
23
24     stats = df.describe()
25     # Escape per cent chars
26     idx = ['count', 'mean', 'std', 'min', '25%', '50%', '75%', 'max']
27     stats = stats.set_index([idx])
28
29     # Save stats as csv
30     stats.to_csv(base+"stats.csv", header=header, float_format="%.3e")
31
32
33 def collect_read_times(file_name):
34     """Read thread times from a file."""
35     f = open(file_name)
36     times = []
37     # Regular expression to find floats
38     time = re.compile("(\\d+\\.\\d+)")
39
40     for line in f:
41         match = time.match(line)
42
43         if (match):
44             t = float(match.group(1))
45             times.append(t)
46
47     return times
48
49 def get_read_times():
50     """Collect timing information about all schedulers in a DataFrame."""
51     schedulers = ["cfq", "deadline", "noop"]
52     base = "../data/"
53     ext = ".log"
54     header=("cfq", "deadline", "noop")
55
56     # Collect all times in one file
57     times = {key: [] for key in schedulers}
58     for s in schedulers:
59         # We have 4 parallel log files for each scheduler
60         for i in [1,2,3,4]:
61             f = base+s+"-"+str(i)+ext
62             times[s].extend(collect_read_times(f))
63     # Return a DataFrame with all timing data
64     df = pd.DataFrame(times)
65     return df
66
67 # Collect statistical data
68 produce_stats()

```

Listing 4: parser.c

```

1 /**

```

```

2  * File: parser.c
3  * Author: Lennart Jern - ens16ljn@cs.umu.se
4  *
5  * This parser extracts search data from command line arguments.
6  *
7  */
8
9  #include <stdio.h>
10 #include <stdlib.h>
11 #include <getopt.h>
12 #include <errno.h>
13 #include <string.h>
14 #include "parser.h"
15
16 /**
17  * parse_arguments - parse command line arguments into SearchData.
18  * The SearchData should be freed by calling SearchData_delete when
19  * you are done.
20  * @param argc -- number of arguments
21  * @param argv -- array of arguments
22  * @return      SearchData containing: dirs to look in, needle to look
23  *              for, file type and number of threads.
24  */
25 SearchData *parse_arguments(int argc, char *argv[]) {
26     SearchData *s_data;
27     char *usage = "Usage: ./mfind [-t{d|f|l}] [-p{nr}thr] [-start1 [start2...]] [-name]";
28
29     // No point to continue if there are less than 3 args
30     if (argc < 3) {
31         fprintf(stderr, "%s\n", usage);
32         exit(EXIT_FAILURE);
33     }
34
35     s_data = malloc(sizeof(SearchData));
36     if (s_data == NULL) {
37         perror("malloc");
38         exit(EXIT_FAILURE);
39     }
40
41     // Set num_threads and type from given arguments
42     if (set_flags(s_data, argc, argv) != 0) {
43         fprintf(stderr, "%s\n", usage);
44         exit(EXIT_FAILURE);
45     }
46     // Make sure the flags are safe
47     check_flags(s_data);
48
49     // Check that there is at least one dir to look in and a name to look for.
50     if (optind >= argc - 1) {
51         fprintf(stderr, "%s\n", usage);
52         free(s_data);
53         exit(EXIT_FAILURE);
54     }
55
56     add_dirs_and_needle(s_data, argv, optind, argc-1);
57
58     return s_data;
59 }
60
61 /**
62  * set_flags - parse the arguments and add corresponding search data
63  * @param s_data -- SearchData to add info to

```

```

64  * @param argc  -- number of arguments
65  * @param argv  -- array of arguments
66  * @return      0 if everything went well, -1 otherwise.
67  */
68  int set_flags(SearchData *s_data, int argc, char *argv[]) {
69      char *optstr = "t:p:";
70      int opt;
71      char type = '\0';
72      int num_threads = 1;
73
74      // Parse flags
75      while ((opt = getopt(argc, argv, optstr)) != -1) {
76          char *end;
77          switch (opt) {
78              case 't':
79                  type = *optarg;
80                  break;
81              case 'p':
82                  errno = 0;
83                  num_threads = strtol(optarg, &end, 10);
84                  if (errno != 0) {
85                      perror("strtol");
86                  }
87                  break;
88              default:
89                  return -1;
90          }
91      }
92
93      s_data->num_threads = num_threads;
94      s_data->type = type;
95      return 0;
96  }
97
98  /**
99   * check_flags - make sure the flags (num_threads and type) are correct and safe
100   * @param s_data -- SearchData to check.
101   */
102  void check_flags(SearchData *s_data) {
103      int num_threads = s_data->num_threads;
104      char type = s_data->type;
105
106      if (num_threads < 1) {
107          fprintf(stderr, "Number of threads must be more than 0!\n");
108          exit(EXIT_FAILURE);
109      } else if (num_threads > MAXTHREADS) {
110          fprintf(stderr, "Too many threads! Resetting to 1.\n");
111          num_threads = 1;
112      }
113      if (type != 'd' && type != 'f' && type != 'l' && type != '\0') {
114          fprintf(stderr, "Type must be d|f|l.\n");
115          exit(EXIT_FAILURE);
116      }
117
118      s_data->num_threads = num_threads;
119      s_data->type = type;
120  }
121
122  /**
123   * add_dirs_and_needle - parse starting dirs and needle from arguments and
124   * add them to the search data.
125   * @param s_data -- search data to add to

```

```

126  * @param argv -- array of arguments
127  * @param first -- index of first directory in argv
128  * @param last -- index of last element in argv. This should be the needle.
129  */
130  void add_dirs_and_needle(SearchData *s_data, char *argv[], int first, int last) {
131      LinkedList *dirs = List_init();
132      char *needle;
133      int num_dirs = last - first;
134
135      // Add all starting directories to the list of dirs.
136      for (int i = 0; i < num_dirs; i++) {
137          char *new_dir = malloc(strlen(argv[first+i])+1);
138          if (new_dir == NULL) {
139              perror("malloc");
140              exit(EXIT_FAILURE);
141          }
142          strcpy(new_dir, argv[first+i]);
143
144          if (List_append(dirs, (void *)new_dir) == 0) {
145              fprintf(stderr, "Could not add path to list.\n");
146              s_data->error++;
147          }
148      }
149
150      // The last arg is the needle/name to search for
151      needle = argv[last];
152      s_data->needle = malloc(strlen(needle)+1);
153      if (s_data->needle == NULL) {
154          perror("malloc");
155          exit(EXIT_FAILURE);
156      }
157
158      // Add everything to s_data.
159      strcpy(s_data->needle, needle);
160      s_data->directories = dirs;
161  }

```

Listing 5: parser.h

```

1  /**
2   * File: parser.h
3   * Author: Lennart Jern - ens16ljn@cs.umu.se
4   *
5   * This is a header file for a parser that extracts search data from command
6   * line arguments.
7   *
8   */
9
10 #include "list.h"
11
12 #define MAXTHREADS (1)
13
14 typedef struct search_data SearchData;
15
16 struct search_data {
17     char *needle;
18     LinkedList *directories;
19     int num_threads;
20     char type;
21     int num_searchers;
22     unsigned int error;

```



```

23 };
24
25 SearchData *parse_arguments(int argc, char *argv[]);
26 int set_flags(SearchData *s_data, int argc, char *argv[]);
27 void check_flags(SearchData *s_data);
28 void add_dirs_and_needle(SearchData *s_data, char *argv[], int first, int last);

```

Listing 6: list.c

```

1  /**
2   * File: list.c
3   * Author: Lennart Jern - ens16ljn@cs.umu.se
4   *
5   * A simple implementation of a linked list.
6   *
7   */
8
9  #include <stdio.h>
10 #include <stdlib.h>
11 #include <string.h>
12 #include "list.h"
13
14 /**
15  * List_init
16  * Create and initialize a LinkedList.
17  *
18  * @return pointer to list
19  */
20 LinkedList *List_init(void) {
21     LinkedList *lst = calloc(1, sizeof(LinkedList));
22     if (lst == NULL) {
23         fprintf(stderr, "Allocation of memory for linked list failed\n");
24         exit(EXIT_FAILURE);
25     }
26     return lst;
27 }
28
29 /**
30  * Appends an element with the specified value to the list
31  * @param l: List to append to
32  * @param value: Pointer to value to be added
33  * @return 1 on success, otherwise 0.
34  */
35 int List_append(LinkedList *l, void *value) {
36
37     Node *new_node = calloc(1, sizeof(Node));
38     if (new_node == NULL) {
39         fprintf(stderr, "Failed to allocate memory for new node.\n");
40         return 0;
41     }
42     new_node->value = value;
43     new_node->next = NULL;
44
45     if (l->first == NULL) {
46         l->first = new_node;
47         return 1;
48     }
49     Node *node_ptr = l->first;
50
51     while (node_ptr->next != NULL) {
52         node_ptr = node_ptr->next;

```

```

53     }
54     node_ptr->next = new_node;
55     return 1;
56 }
57
58 /**
59  * List_get - get the value of the first node in the list and delete the node
60  * from the list.
61  * @param l -- the list
62  * @return -- pointer to value, remember to free it later.
63  */
64 void *List_get(LinkedList *l) {
65     if (l->first == NULL) {
66         return NULL;
67     }
68     Node *node = l->first;
69     l->first = l->first->next;
70     void *value = node->value;
71     free(node);
72     node = NULL;
73     return value;
74 }
75
76 /**
77  * List_sort
78  * Sort the list lst by selection sort, using the comparison function comp.
79  * @param lst: LinkedList to sort
80  * @param comp: function used to compare the values of two nodes to determine
81  *              what order they should be placed in.
82  */
83 void List_sort(LinkedList *lst, int (*comp)(void *value1, void *value2)) {
84     if (lst->first == NULL) {
85         // Empty list, nothing to do
86         return;
87     }
88     Node *boundary = lst->first; // ordered nodes before this
89     Node *smallest = lst->first; // should be placed next in order
90     Node *last_sorted = NULL; // add smallest after this one
91     Node *current = lst->first;
92     Node *prev = NULL;
93     Node *before_smallest = NULL;
94
95     // run untill the whole list is sorted
96     while (boundary->next != NULL) {
97         smallest = boundary;
98         current = boundary;
99         prev = NULL;
100        before_smallest = NULL;
101        // loop through the unordered part of the list and pick out the node
102        // with the "smallest" value
103        while (current->next != NULL) {
104            prev = current;
105            current = current->next;
106            if (comp(smallest->value, current->value) < 0) {
107                smallest = current;
108                before_smallest = prev;
109            }
110        }
111
112        // Do we have to move smallest?
113        if (before_smallest != NULL) {
114            // cut out smallest

```

```

115         before_smallest->next = smallest->next;
116         if (last_sorted == NULL) {
117             // place it first
118             smallest->next = lst->first;
119             lst->first = smallest;
120         } else {
121             // place it after last_sorted
122             smallest->next = last_sorted->next;
123             last_sorted->next = smallest;
124         }
125     }
126     // update last_sorted and boundary
127     last_sorted = smallest;
128     boundary = smallest->next;
129 }
130 }
131
132 /**
133  * List_remove
134  * Removes an element from the bottom of the list and frees the allocated memory
135  * by calling delete_value(value).
136  *
137  * @param lst: the list to remove from
138  * @return 1 on successful removal, 0 if no node was removed.
139  */
140 int List_remove(LinkedList *lst, void (*delete_value)(void *value)) {
141     Node *node = lst->first;
142     Node *prev = node;
143     if (node == NULL) {
144         return 0;
145     }
146     if (node->next == NULL) { // Only one element in list
147         delete_value(node->value);
148         free(node);
149         lst->first = NULL;
150         return 1;
151     }
152     while (node->next != NULL) {
153         prev = node;
154         node = node->next;
155     }
156     prev->next = NULL;
157     delete_value(node->value);
158     free(node);
159     return 1;
160 }
161
162 /**
163  * List_delete
164  * Frees all memory allocated by the nodes in the LinkedList
165  * using delete_value(value) and frees the list itself after that.
166  *
167  * @param lst: LinkedList to free.
168  * @param delete_value: function used to free the memory allocated by the value
169  * of a node.
170  */
171 void List_delete(LinkedList *lst, void (*delete_value)(void *value)) {
172     if (lst) {
173         while(List_remove(lst, delete_value)) ;
174         free(lst);
175     }
176 }

```

```

177
178 /**
179  * List_print
180  * Prints the LinkedList using the function provided.
181  *
182  * @param lst:    LinkedList to print.
183  * @param print:  function used to print the node values.
184  */
185 void List_print(LinkedList *lst, void (*print)(void *value)) {
186     Node *current = lst->first;
187
188     while (current != NULL) {
189         print(current->value);
190         current = current->next;
191     }
192 }

```

Listing 7: list.h

```

1 /**
2  * File: list.h
3  * Author: Lennart Jern - ens16ljn@cs.umu.se
4  *
5  * This is a header file for my own implementation of a linked list.
6  *
7  */
8
9 typedef struct linked_list LinkedList;
10 typedef struct node Node;
11
12 struct node {
13     void *value;
14     struct node *next;
15 };
16
17 struct linked_list {
18     Node *first;
19 };
20
21 struct user_info {
22     unsigned int uid;
23     char *uname;
24 };
25
26 /**
27  * List_init
28  * Create and initialize a LinkedList.
29  *
30  * @return pointer to list
31  */
32 LinkedList * List_init(void);
33
34 /**
35  * Appends an element with the specified value to the list
36  * @param l:    List to append to
37  * @param value: Pointer to value to be added
38  * @return      1 on success, otherwise 0.
39  */
40 int List_append(LinkedList *l, void *value);
41
42 /**

```

```

43  * List_get - get the first element from the list. The element is removed
44  * from the list so remember to free it when you are done.
45  * @param l -- the list
46  * @return -- a pointer to the first value in the list
47  */
48  void *List_get(LinkedList *l);
49
50  /**
51  * List_sort
52  * Sort the list lst by selection sort, using the comparison function comp.
53  * @param lst: LinkedList to sort
54  * @param comp: function used to compare the values of two nodes to determine
55  *               what order they should be placed in.
56  */
57  void List_sort(LinkedList *lst, int (*comp)(void *value1, void *value2));
58
59  /**
60  * List_remove
61  * Removes an element from the bottom of the list and frees the allocated memory
62  * by calling delete_value(value).
63  *
64  * @param lst: the list to remove from
65  * @return 1 on successful removal, 0 if no node was removed.
66  */
67  int List_remove(LinkedList *lst, void (*delete_value)(void *value));
68
69  /**
70  * List_delete
71  * Frees all memory allocated by the nodes in the LinkedList
72  * using delete_value(value) and frees the list itself after that.
73  *
74  * @param lst: LinkedList to free.
75  * @param delete_value: function used to free the memory allocated by the value
76  *                       of a node.
77  */
78  void List_delete(LinkedList *lst, void (*delete_value)(void *value));
79
80  /**
81  * List_print
82  * Prints the LinkedList using the function provided.
83  *
84  * @param lst: LinkedList to print.
85  * @param print: function used to print the node values.
86  */
87  void List_print(LinkedList *lst, void (*print)(void *value));

```

## Listing 8: Makefile

```

1  SOURCE=mfind.c list.c
2  OBJECTS=mfind.o list.o parser.o
3  FLAGS=-std=c11 -Wall -pedantic -Werror -pthread
4
5  all: $(OBJECTS)
6      gcc $(OBJECTS) -pthread -o mfind
7      gcc $(FLAGS) crazy_search.c -o crazy_search
8
9  mfind.o: mfind.c
10     gcc $(FLAGS) -c mfind.c
11
12  list.o: list.c list.h
13     gcc $(FLAGS) -c list.c

```

```

14
15 parser.o: parser.c parser.h list.h
16     gcc $(FLAGS) -c parser.c
17
18 debug: FLAGS+=-DDEBUG -g
19 debug: all
20
21 test: all
22     ./mfind -td -p2 . .. fail mfind
23
24 time: FLAGS+=-DTIME
25 time: all
26
27 memtest: all
28     valgrind ./mfind -td -p2 . .. mfind
29
30 clean:
31     rm -f mfind *.o

```

Listing 9: generate\_test\_files.sh

```

1  #!/bin/bash
2
3  # File: generate_test_files.sh
4  # Author: Lennart Jern - ens16ljn@cs.umu.se
5  #
6  # Generate a file tree to do tests on.
7
8  # Starting directory
9  # START=/run/media/lennart/KINGSTON
10 START=/run/media/lennart/Verbatim
11
12 # Move to correct directory/device
13 cd $START
14
15 # Remove tree if existent
16 rm -r test_files
17
18 # Create directory to hold all test files
19 mkdir test_files
20 cd test_files
21
22 # Split up the files between a few directories
23 for DIR in a b c d e; do
24     mkdir $DIR
25     cd $DIR
26     # Create empty files
27     for F in $(seq 1 20); do
28         touch $F
29     done
30
31     for D in f g h i j; do
32         mkdir $D
33         cd $D
34     done
35
36     cd "$START/test_files"
37 done
38
39 # Big files
40 mkdir bigs

```

```
41 cd bigs
42
43 # Generate files with lots of zeros...
44 for F in $(seq 1 5); do
45     head -c 50M < /dev/zero > "file$F"
46 done
47
48 cd "$START/test_files"
49 # Generate deep folders
50 for i in $(seq 1 10); do
51     mkdir "deep$i"
52     cd "deep$i"
53     for D in $(seq 1 100); do
54         mkdir "dir$D"
55         cd "dir$D"
56     done
57     cd "$START/test_files"
58 done
```